



HIETANEN CONSULTANCY • CUSTOMER BRIEFING • SEPTEMBER 2026



Fuuga

The sovereign LLM stack
for .NET.

Train, fine-tune, compress and serve your own language models on .NET — no Python, and no data leaving the building.

Enterprises rent their AI, and cannot maintain the AI they build.



Rented models, no control

Data leaves the company, every request is billed, and the vendor can change or retire the model under a live product.



In-house models today are a specialist's side project

A data scientist builds them in Python tools the engineering team does not use. When that person leaves, nobody can maintain them.




Regulation asks who owns the model

For a high-risk use such as credit scoring, the EU AI Act requires a technical file describing the training data. DORA requires an exit plan for each critical ICT supplier.

THE IDEA

Ahead where your business is.



Retrained on your work

Your decisions, formats and code conventions, taught to an open model.

Grounded in all your documents

Searched in-house, including the ones that may never leave.

Measured both ways

The gain and the small loss, on your own held-out tasks, before release.

Your own AI model, built and maintained by the engineers you already have.

The complete toolchain to make a language model, or retrain an existing one — on .NET, the platform banks and insurers already run.

- Starts from the best open models and teaches them your data
- Runs on ordinary GPUs or CPU servers you already own
- Speaks the OpenAI and Anthropic interfaces; nothing is rewritten

for the technical reader — the common path

```
> fuuga scaffold finetune  
> fuuga transfer --recipe my.json  
> fuuga sft --recipe my.json  
> fuuga eval --recipe my.json  
> fuuga export gguf --tokenizer tok/  
> fuuga-serve --checkpoint best/
```

Start from an open model. Make it yours. Keep it in-house.

01

Ingest

Company data becomes a traceable corpus.

02

Transfer

Start from nine open model families.

03

Fine-tune

Teach it the company's work on ordinary GPUs.

04

Compress

Shrink it to run on servers already owned.

05

Serve

OpenAI- and Anthropic-compatible, in-house.

Every step is measured, so an agent can run the loop — and only keeps a model that is better.

Not a roadmap. A working stack in the repository.

Area	In the codebase now	Status
Training	Modern transformer architectures, models larger than the GPU	Shipped
Fine-tuning	SFT, DPO, GRPO, LoRA and QLoRA, typed judge	Shipped
Model operations	Transfer, distillation, merging, quantisation, export	Shipped
Serving	OpenAI and Anthropic APIs, batching, guardrails, MCP	Shipped
Metered hosting	Per-token pricing and prepaid credit per API key	Shipped
Agents	Sandboxed tools, cost-aware routing, knowledge-graph RAG	Shipped
Quality	Unit and property tests across every module	Shipped

Open models caught up and regulation arrived — yet Microsoft's .NET path is still a rented one.

Open models are good enough

Adapted to a company's own cases, open weights beat a generic hosted model.

Hosted prices must rise

Today's per-token prices are subsidised by investor capital.

Sovereignty is required

Credit scoring becomes high-risk under the EU AI Act in December 2027; DORA has applied to finance since January 2025.

.NET's path is rented

Microsoft's route is Copilot and Azure OpenAI: OpenAI and Anthropic models, hosted, not owned.

WHERE IT PAYS FIRST

Typed decisions: the part of AI a business can actually measure.

The data they have

Outcomes, not a corpus

Every company records outcomes — approved, declined, fraud. Those are the training labels.

The number they get

An answer to "is it better?"

A decision has ground truth, so accuracy and calibration are measured on the customer's own cases.

The line it moves

Act, or escalate

Confident cases are handled by software; the rest go to a person. The share automated is the return.

Nothing here generates text. One pass returns a calibrated probability for each option — logged, audited and priced.

The bill is the business case. Move a workload only when the numbers say so.

01

Connect

Same OpenAI and Anthropic interfaces; point the app at a new address.

02

Route

Sensitive and routine requests in-house, the rest to the subscription.

03

Learn

Train the in-house model on the company's own traffic.

04

Switch

When the measured comparison passes, the subscription shrinks.

A rising per-request bill becomes a fixed cost — and the weights leave in standard formats, so it runs anywhere, including without us.

Hietanen Consultancy • Fuuga • 09 • `fuuga-serve --api-compat both; cost-aware router; distill --teacher-endpoint; export gguf / onnx`

Today's options are a Python team or a rented model. Fuuga adds a third.

What the buyer gets	Fuuga	Python stack	.NET libraries	Copilot and hosted APIs
Train and fine-tune a language model	Yes	Yes	Small classifiers	On their servers
Data stays on your hardware	Yes	Yes	Inference only	No
Maintained by the .NET team	Yes	New team	Yes	Yes
Training provenance	Built in	Assemble it	—	Vendor statement
You own the weights	Yes	Yes	Import only	No
Cost as usage grows	Fixed	Fixed	Fixed	Per request

Python: PyTorch, HuggingFace, vLLM · .NET libraries: Semantic Kernel, LLamaSharp, ML.NET · Hosted: Copilot, Azure OpenAI, Mistral

One engineer, AI-assisted, shipped the stack. The company grows as customers sign.



Tuomas Hietanen — Founder and CEO

Nine years as CTO of two regulated lenders, UK and US. Maintainer of SQLProvider and LINQKit.

Who owns what	Today	Next hire
Platform	Founder, AI-assisted	Engineers
ML delivery	Founder and network	ML engineers
Enterprise sales	Founder	Head of sales
Compliance	Founder	Compliance lead

Identified, and named once agreements are in writing: a second senior team member, and partners for data, compliance and ML advisory. The network is senior engineers we have shipped with, joining on contract. The founder hands over in the order shown, platform first, as signed revenue pays for each hire; public code, full tests and written recipes let them run the loop from the first day.

Start with a pilot on outcomes you already record.

01 · Pilot

Fixed scope

A calibrated decision engine in production on your own outcomes.

02 · We run it

In your environment

Retraining, drift monitoring, evidence for your supervisor, on-call.

03 · You run it

Supported builds

Signed releases, an SLA, provenance — and onboarding each new engineer.

04 · Leave any time

Nothing locked in

Public-domain code, standard weight formats, every recipe written down.

Tuomas Hietanen · hello@hietanen.co.uk · hietanen.co.uk/projects/fuuga/